

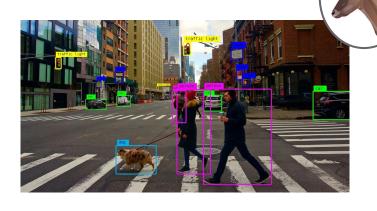
OpenCV Conceptual Architecture eLand

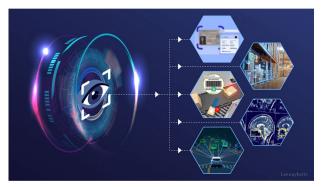
Eric Rodrigues, Laura Marin, Alp Baran Sirek, Negar Khalilazar, Danny Le



What Is Computer Vision?

- Computer vision is a tool that enables computers to interpret visual information from the world
- We use it when cars track objects in real time, in augmented reality tools, etc
- Although the potential of computer vision is exciting, it also presents several challenges







Enter OpenCV

- OpenCV is one of the most widely used tools for addressing challenges of computer vision
- It is an open-source library designed for real-time image processing and computer vision applications
- The library, however, has expanded significantly now supporting additional languages and tools







Today You Will Learn



- The Conceptual architecture of OpenCV
- How the system is organized and how its parts interact
- The illustration of the system's structure, design and flow of control and data
- Division of Responsibilities and work flow
- Practical applications





What the System Does



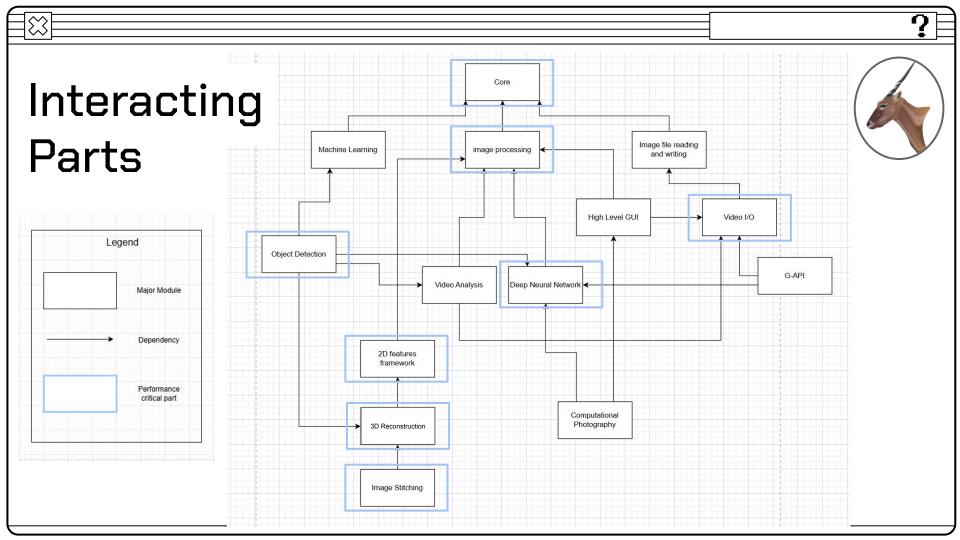
Main Modules

- Core functionality
- Image Processing
 - Image File I/O
 - Video I/O
 - High-level GUI
- Video Analysis
- 3D Reconstruction
- 2D Features Framework
- Object Detection
- Deep Neural Network module
- Machine Learning
- G-APIComputational Photography
- Images stitching

Architecture Style

- Pipe and Filter?
- Repository?
- Modular and Layered

									?
 -	łię	jh-L		Arc		ect	ure Dia	agrar	n
Object Detection	Deep Neural Network		Machine L	Learning	Image Stitching		G-API		
Layer 3: Analysis and Feature Extraction									
Video Analysis		2D Features Framework			3D R	econstruction		Legend	
Layer 2: Base Processing and I/O									Layer
Image Processing	Image Processing Vio		o I/O High		h level GUI C		mputational Photography		Major Module
Layer 1: Core									
Core Functionality									
			1						

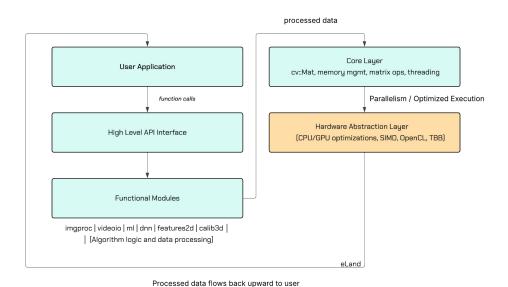




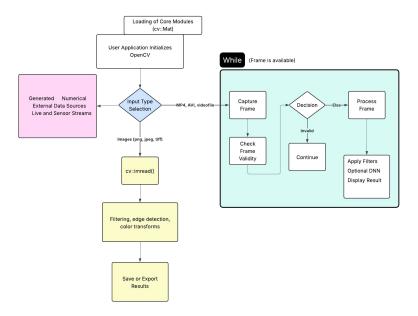
Control Flow

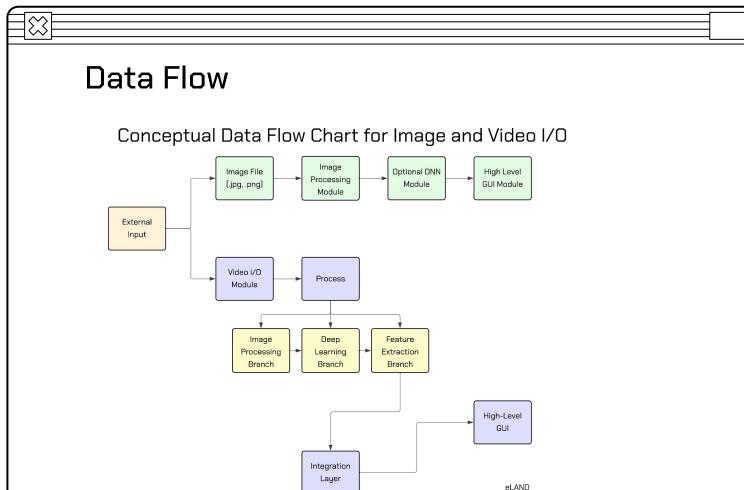


Conceptual-General Control Flow



A More Specific Control Flow

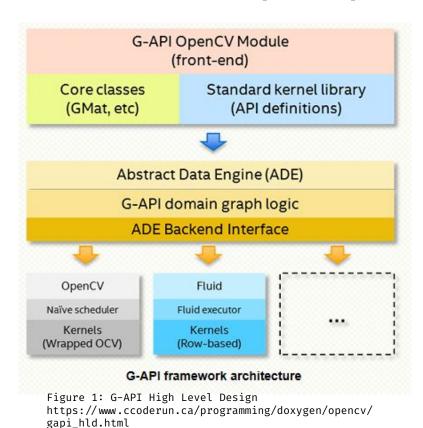








Control Flow (G-API)



G-API High-level design

- Defines processing as a graph of operations instead of step-by-step calls
- API Layer: user builds pipelines using GMat, GScalar, etc.
- Graph Compiler: optimizes and schedules execution
- Backends: choose best hardware automatically (CPU, GPU, OpenCL, etc.)
- Enables parallelism and zero-copy data flow between node



Concurrency



- Fully re-entrant: functions safe to call from multiple threads.
- Thread-safe data: cv::Mat uses atomic reference counting.
- Built-in parallelism: uses TBB / OpenMP internally.
- Custom threading: via cv::parallel_for_() and setNumThreads().
- Device-level concurrency: CPU + GPU tasks can run in parallel.

<u>efficient</u>, <u>scalable real-time processing</u>

Initial OpenCV Release

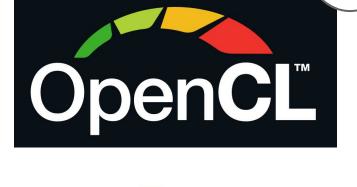
- OpenCV's early versions were primarily optimized for CPU-based image processing
- Capable of low level operations such as filtering, edge detection, and feature extraction
- Supported C and C++ ONLY

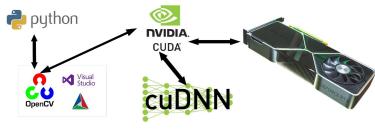




Performance Enhancements

- Around version 2.0 the CUDA module was introduced to leverage NVIDIA GPUs for parallel computation
- OpenCV later integrated OpenCL support through the Transparent API (T-API), which allowed hardware-agnostic acceleration





Expansion of Bindings

- OpenCV introduced bindings for additional languages such as Python, Java, and JavaScript
- Extended its reach to mobile platforms including Android and iOS



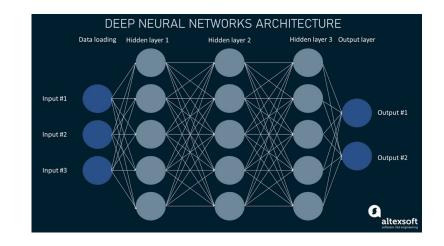




Deep Neural Network



- OpenCV 3.3 (released in 2017), added the Deep Neural Network (DNN) module
- This enabled users to import and run pre-trained neural networks from frameworks
- This addition did not disrupt or modify the pre-existing modules preserving backward compatibility





Modern Al Applications

- Today, OpenCV plays a central role in the development of artificial intelligence systems
- Often used to prepare, augment, and transform data for training deep learning models
- Examples seen are autonomous driving and facial recognition

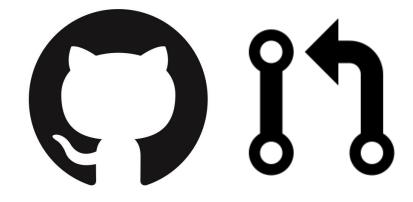




Division of Responsibilities



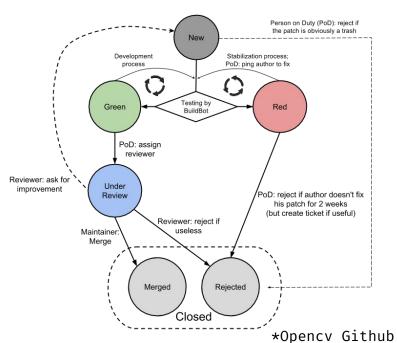
- Users → contributors →
 committers → release managers
 (→ technical group)
- Changes done via pull requests, review and test
- Decisions publicly recorded and released



How Work Flows



- Libraries are split by function
- QA: style guide, tests, cross-platform builds
- Decisions publicly recorded and released

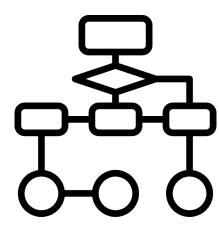




Division of Responsibilities



- Each step in the pipeline has an owner
- Pull requests must define the API, add tests, document, show performance impact
- Branch policy: small fixes to current stable and big features next major





Practical Applications (1.0)

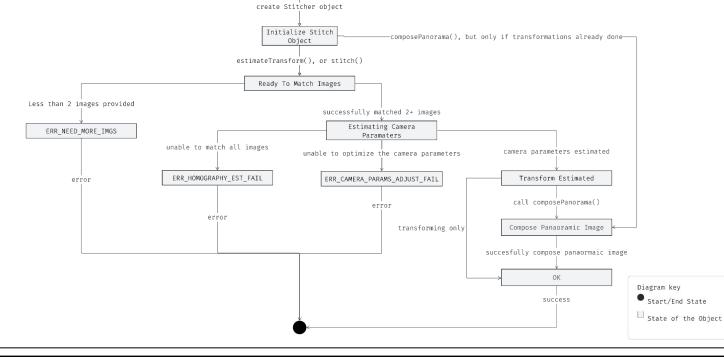
- OpenCV supports many situations where a computer needs to process and analyze visual information. One such case is the planetary rovers and drones.
 - Planetary rovers and drones require maximum situational awareness, as failure can be costly.
 - Multiple cameras require multiple video feeds to review, and can create overlapping fields of view.
 - A rotating camera cannot view all angles at once, and is extra complexity.

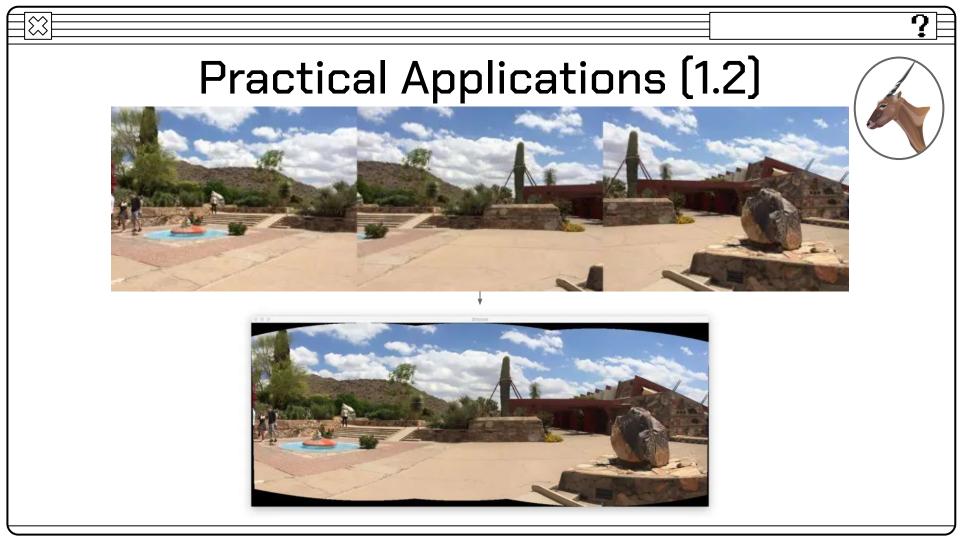


Practical Applications (1.1)

Stitcher class (homography model):
 Transforming multiple images/videos into one panoramic image/video.







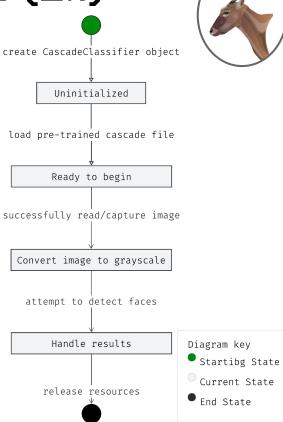
Practical Applications (2.0)

- Another case is facial recognition.
 - Easy authentication
 - Observing animals
 - Human-Robot interaction
 - Personal and Private security
 - Automated safety precautions
 - Forensic Analysis



Practical Applications (2.1)

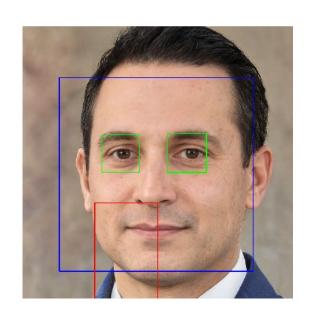
 CascadeClassifier class: scans images with multi-stage facial detection to efficiently reject non-facial regions early and focusing computational effort only on promising areas

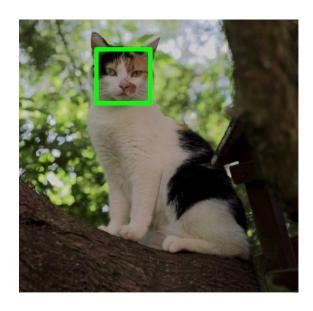


Practical Applications (2.2)

- Capable of distinguishing people from animals







Lessons Learned - Begin the analysis earlier - Start Creating Diagrams

- Better documentation of our thought process

earlier

